Lecture 3: Theoretical Analysis of Boosting Methods

Tianpei Xie

Feb. 6th., 2023

Contents

Algorithm 1.1
The boosting algorithm AdaBoost

Given: $(x_1, y_1), \ldots, (x_m, y_m)$ where $x_i \in \mathcal{X}$, $y_i \in \{-1, +1\}$.
Initialize: $D_1(i) = 1/m$ for $i = 1, \ldots, m$.
For $t = 1, \ldots, T$:

- Train weak learner using distribution D_t.
- Get weak hypothesis $h_t : \mathcal{X} \to \{-1, +1\}$.
- Aim: select h_t to minimalize the weighted error:

$$\epsilon_t \doteq \mathbf{Pr}_{i \sim D_t}[h_t(x_i) \neq y_i].$$

- Choose $\alpha_t = \dfrac{1}{2} \ln\left(\dfrac{1-\epsilon_t}{\epsilon_t}\right)$.
- Update, for $i = 1, \ldots, m$:

$$D_{t+1}(i) = \frac{D_t(i)}{Z_t} \times \begin{cases} e^{-\alpha_t} & \text{if } h_t(x_i) = y_i \\ e^{\alpha_t} & \text{if } h_t(x_i) \neq y_i \end{cases}$$

$$= \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t},$$

where Z_t is a normalization factor (chosen so that D_{t+1} will be a distribution).

Output the final hypothesis:

$$H(x) = \text{sign}\left(\sum_{t=1}^{T} \alpha_t h_t(x)\right).$$

Figure 1: AdaBoost Algorithm [Schapire and Freund, 2012]

1 Boosting Algorithm

1.1 AdaBoost

-

1.2 Functional Gradient Descent

-

1.3 Gradient Boost

-

Algorithm 7.3
AnyBoost, a generic functional gradient descent algorithm

Goal: minimization of $\mathcal{L}(F)$.

Initialize: $F_0 \equiv 0$.

For $t = 1, \ldots, T$:

- Select $h_t \in \mathcal{H}$ that maximizes $-\nabla \mathcal{L}(F_{t-1}) \cdot h_t$.
- Choose $\alpha_t > 0$.
- Update: $F_t = F_{t-1} + \alpha_t h_t$.

Output F_T.

Figure 2: Gradient Boost Tree Algorithm [Hastie et al., 2009]

Algorithm 10.3 *Gradient Tree Boosting Algorithm.*

1. Initialize $f_0(x) = \arg\min_\gamma \sum_{i=1}^{N} L(y_i, \gamma)$.

2. For $m = 1$ to M:

    (a) For $i = 1, 2, \ldots, N$ compute

    $$r_{im} = -\left[\frac{\partial L(y_i, f(x_i))}{\partial f(x_i)}\right]_{f=f_{m-1}}.$$

    (b) Fit a regression tree to the targets r_{im} giving terminal regions R_{jm}, $j = 1, 2, \ldots, J_m$.

    (c) For $j = 1, 2, \ldots, J_m$ compute

    $$\gamma_{jm} = \arg\min_\gamma \sum_{x_i \in R_{jm}} L\left(y_i, f_{m-1}(x_i) + \gamma\right).$$

    (d) Update $f_m(x) = f_{m-1}(x) + \sum_{j=1}^{J_m} \gamma_{jm} I(x \in R_{jm})$.

3. Output $\hat{f}(x) = f_M(x)$.

Figure 3: Gradient Boost Tree Algorithm [Hastie et al., 2009]

2 Theoretical Guarantee for Boosting

- **Remark** (**Data**)
  Define an **observation** as a d-dimensional vector x. The *unknown* nature of the observation is called a **class**, denoted as y. The domain of observation is called an **input space** *or* **feature space**, denoted as $\mathcal{X} \subset \mathbb{R}^d$, whereas the domain of class is called the **target space**, denoted as $\mathcal{Y}$. For **classification task**, $\mathcal{Y} = \{1, \ldots, M\}$; and for **regression task**, $\mathcal{Y} = \mathbb{R}$. A **concept** $c : \mathcal{X} \to \mathcal{Y}$ is the *input-output association* from the nature and is *to be learned* by **a learning algorithm**. Denote $\mathcal{C}$ as *the set of all concepts* we wish to learn as the **concept class**. The learner is requested to output a *prediction rule*, $h : \mathcal{X} \to \mathcal{Y}$. This function is also called a **predictor**, a **hypothesis**, or a **classifier**. The predictor can be used to predict the label of new domain points. Denote a collection of n **samples** as

$$\mathcal{D} \equiv \mathcal{D}_n = ((X_1, Y_1), \ldots, (X_n, Y_n)) \equiv ((X_1, c(X_1)), \ldots, (X_n, c(X_n))).$$

  Note that $\mathcal{D}_n$ is a finite **sub-sequence** in $(\mathcal{X} \times \mathcal{Y})^n$.

- **Definition** (**Generalization Error in Deterministic Scenario**) [Mohri et al., 2018]
  Under *a deterministic scenario*, **generalization error** or the **risk** or simply **error** for the *classifier* $h \in \mathcal{H}$ is defined as

$$L(h) \equiv L_{\mathcal{P},c}(h) = \mathcal{P}\{h(X) \neq c(X)\} \equiv \mathbb{E}_X\left[\mathbb{1}\{h(X) \neq c(X)\}\right] \tag{1}$$

  with respect to the concept $c \in \mathcal{C}$ and *the feature distribution* $\mathcal{P} \equiv \mathcal{P}_X$.

- **Definition** (**Empirical Error or Training Error**)
  Given the data $\mathcal{D}$, the **training error** or the **empirical error/risk** of a hypothesis $h \in \mathcal{H}$ is defined as

$$\widehat{L}(h) \equiv \widehat{L}_{\mathcal{D}}(h) = \frac{1}{n}\sum_{i=1}^{n}\mathbb{1}\{h(X_i) \neq Y_i\} = \frac{1}{n}|\{i : h(X_i) \neq Y_i\}| := \widehat{\mathbb{E}}\left[\mathbb{1}\{h(X) \neq Y\}\right]$$

  where either $Y = c(X)$ or Y is a random variable associated with X.

- **Definition** (**The Realizable Assumption**)
  There exists $h^* \in \mathcal{H}$ s.t. $L_{\mathcal{P},c}(h^*) = 0$.

- **Definition** (**PAC Learnability**)
  A hypothesis class $\mathcal{H}$ is **PAC learnable** if there exist a function $m_{\mathcal{H}} : (0,1)^2 \to \mathbb{N}$ and a learning algorithm with the following property: For every $\epsilon, \delta \in (0,1)$, *for every distribution* $\mathcal{P}$ *over* $\mathcal{X}$, and *for every labeling function* $c : \mathcal{X} \to \{0,1\}$, if *the realizable assumption* holds with respect to $\mathcal{H}$, $\mathcal{P}$, c, then when running the learning algorithm on $m \geq m_{\mathcal{H}}(\epsilon, \delta)$ i.i.d. examples generated by $\mathcal{P}$ and labeled by c, the algorithm returns a hypothesis h such that, with probability of at least $1 - \delta$ (over the choice of the examples),

$$L_{\mathcal{P},c}(h) \leq \epsilon.$$

2.1 Weak Learner

- **Definition** (γ-**Weak Learnability**) [Schapire and Freund, 2012, Shalev-Shwartz and Ben-David, 2014]

4

A learning algorithm, $\mathcal{A}$, is a γ-**weak-learner** for a class $\mathcal{H}$ if there exists a function $m_{\mathcal{H}}$: $(0,1) \to \mathbb{N}$ such that for **every** $\delta \in (0,1)$, **for every distribution** $\mathcal{P}$ over $\mathcal{X}$, and **for every labeling function** $c : \mathcal{X} \to \{-1,+1\}$, if *the realizable assumption holds* with respect to $\mathcal{H}$, $\mathcal{P}$, c, then when running the learning algorithm on $m \geq m_{\mathcal{H}}(\delta)$ i.i.d. examples generated by $\mathcal{P}$ and labeled by c, the algorithm returns a hypothesis h such that, with probability of at least $1 - \delta$,

$$L_{\mathcal{P},c}(h) \leq \frac{1}{2} - \gamma.$$

A hypothesis class $\mathcal{H}$ is γ-**weak-learnable** if there exists a γ-*weak-learner* for that class.

- **Remark** We call PAC learnalbe **the strong learnable**.

- **Remark** (**Weak Learner Without Accuracy Guarantee**)
  Unlike *the PAC learner*, who guarantees that *with high probability* the generalization error rate is less than ϵ **for all** ϵ, a γ-**weak-learner** guarantees that with high probability, the error rate is less than ϵ **for some** $\epsilon = 1/2 - \gamma$, i.e. *less than half with a margin γ*.

  In other word, *under the realizablity assumption*, it is expected that **with more data**, **a PAC learner** can learn the "*true*" labeling function behind the data, (i.e. **zero generalization error** with high probability). While a γ-**weak-learner** can only get **slightly better than random guess** and it is **not** *expected to have* **lower error rate** even if more data are available.

- **Remark** (**Weak Learner is as Hard as PAC Learner**)
  The fundamental theorem of learning states that if a hypothesis class $\mathcal{H}$ has a VC dimension d, then *the sample complexity of PAC learning $\mathcal{H}$ satisfies $m_{\mathcal{H}}(\epsilon, \delta) \geq C_1(d + \log(1/\delta))/\epsilon$*, where C_1 is a constant. Applying this with $\epsilon = 1/2 - \gamma$ we immediately obtain that **if $d = \infty$ then $\mathcal{H}$ is not γ-weak-learnable**.

  This implies that from **the statistical perspective** (i.e., if we ignore *computational complexity*), **weak learnability** *is also characterized by the VC dimension of $\mathcal{H}$ and therefore is just* **as hard as PAC (strong) learning**. However, when we do consider **computational complexity**, the potential advantage of weak learning is that maybe there is *an algorithm that satisfies the requirements of weak learning and* **can be implemented efficiently**.

2.2 Training Error Bounds

- **Remark** Recall that $h_t \in \mathcal{H}$ are base learners for $t \in [1, T]$, and $(\alpha_1, \ldots, \alpha_T) \in \Sigma_T$. The *combined learner* is

$$H(x) := \operatorname{sgn}\left(\sum_{t=1}^{T} \alpha_t h_t(x)\right)$$

- The space of all such combined classifiers is defined as below:

  Definition (**Class of Linear Combination of Base Hypothese**)
  Define the class of T **linear combinations of base hypotheses** from $\mathcal{H}$ as

$$L(\mathcal{H}, T) := \left\{ \operatorname{sgn}\left(\sum_{t=1}^{T} \alpha_t h_t(\cdot)\right) : \alpha \in \mathbb{R}^T, h_t \in \mathcal{H}, t = 1, \ldots, T \right\} \tag{2}$$

5

- **Definition** Define Σ_n as the space of all **linear threshold functions**

$$\Sigma_n := \{\text{sgn}\,(\langle w,\,x \rangle) : w \in \mathbb{R}^n\}\,.$$

Thus $L(\mathcal{H},T) = \{\sigma\,(h_1(x),\ldots,h_T(x)) : \sigma \in \Sigma_T\}$

- **Proposition 2.1** *(Training Error Bound for AdaBoost) [Schapire and Freund, 2012]*
  Given the notation of Adaboost algorithm, let $\gamma_t = 1/2 - \epsilon_t$, and let $\mathcal{D}_1$ be an arbitrary initial distribution over the training set. Then **the weighted training error** *of the combined classifier $\mathcal{H}$ with respect to $\mathcal{D}_1$ is bounded as*

$$\widehat{L}_{\mathcal{D}_1}(H) \leq \prod_{t=1}^{T} \sqrt{1 - 4\gamma_t^2} \leq \exp\left(-2\sum_{t=1}^{T}\gamma_t^2\right). \tag{3}$$

2.3 Generalization Error Bounds for Finite Hypothesis Class

- **Definition** (**Restriction of $\mathcal{H}$ to $\mathcal{D}$**).
  Let $\mathcal{H}$ be a class of functions from $\mathcal{X}$ to $\{0,1\}$ and let $\mathcal{D} = \{x_1,\ldots,x_m\} \subset \mathcal{X}$.

  The restriction of $\mathcal{H}$ to $\mathcal{D}$ is *the set of functions* from $\mathcal{D}$ to $\{0,1\}$ that can be *derived from $\mathcal{H}$*. That is,

$$\mathcal{H}_{\mathcal{D}} := \{(h(x_1),\ldots,h(x_m)) : h \in \mathcal{H}\}\,,$$

  where we **represent** each function from $\mathcal{X}$ to $\{0,1\}$ as a **vector** in $\{0,1\}^{|\mathcal{D}|}$.

- **Definition** (**Shattering**).
  A hypothesis class $\mathcal{H}$ **shatters** a finite set $\mathcal{D} \subset \mathcal{X}$ if **the restriction of $\mathcal{H}$ to $\mathcal{D}$** is the set of **all functions** from $\mathcal{D}$ to $\{0,1\}$. That is,

$$|\mathcal{H}_{\mathcal{D}}| = 2^{|\mathcal{D}|}.$$

- **Definition** (**Growth Function**).
  Let $\mathcal{H}$ be a hypothesis class. Then **the growth function of $\mathcal{H}$**, denoted $\tau_{\mathcal{H}} : \mathbb{N} \to \mathcal{N}$, is defined as

$$\tau_{\mathcal{H}}(m) := \max_{\mathcal{D} \subset \mathcal{X}:|\mathcal{D}|=m} |\mathcal{H}_{\mathcal{D}}|\,.$$

  In words, $\tau_{\mathcal{H}}(m)$ is **the number of different functions** from a set $\mathcal{D}$ of **size** m to $\{0,1\}$ that can be obtained by **restricting $\mathcal{H}$ to $\mathcal{D}$**.

- **Lemma 2.2** *(Sauer's Lemma). [Shalev-Shwartz and Ben-David, 2014, Mohri et al., 2018]*
  Let $\mathcal{H}$ be a hypothesis class with $VCdim(\mathcal{H}) \leq d < \infty$. Then, for all $m \geq d+1$,

$$\tau_{\mathcal{H}}(m) \leq \sum_{i=0}^{d} \binom{m}{i} \leq \left(\frac{em}{d}\right)^d. \tag{4}$$

- **Proposition 2.3** *(**Generalization Bound via Growth Function**) [Mohri et al., 2018]*
  Let $\mathcal{H}$ be a family of functions taking values in $\{-1, +1\}$. Then, for any $\delta > 0$, with probability at least $1 - \delta$, for any $h \in \mathcal{H}$,

$$L(h) \leq \widehat{L}_m(h) + \sqrt{\frac{2 \log \tau_{\mathcal{H}}(m)}{m}} + \sqrt{\frac{\log(1/\delta)}{2m}} \tag{5}$$

  Growth function bounds can be also derived directly (without using Rademacher complexity bounds first). The resulting bound is then the following:

$$\mathcal{P}\left\{\exists h \in \mathcal{H}, \left|L(h) - \widehat{L}_m(h)\right| > \epsilon\right\} \leq 4\tau_{\mathcal{H}}(2m) \exp\left(-\frac{m\epsilon^2}{8}\right) \tag{6}$$

  which only differs from (5) by constants.

- The following lemma shows that the VC dimension of Σ_T is T.

  Lemma 2.4 *[Schapire and Freund, 2012]*
  *The space Σ_n of **linear threshold functions** over $\mathbb{R}^n$ has **VC-dimension** n.*

- **Lemma 2.5** *(**Growth Number of Combined Hypothesis Class, Finite Hypothesis Class**) [Schapire and Freund, 2012, Shalev-Shwartz and Ben-David, 2014]*
  *Assume $\mathcal{H}$ is **finite**. Let $m \geq T \geq 1$. For any set $\mathcal{D}$ of m points, the number of dichotomies realizable by $L(\mathcal{H}, T)$ is bounded as follows:*

$$|L(\mathcal{H}, T)| \leq \tau_{L(\mathcal{H}, T)}(m) \leq \left(\frac{em}{T}\right)^T |\mathcal{H}|^T. \tag{7}$$

- **Theorem 2.6** *(**Generalization Bound for AdaBoost, Finite Hypothesis**) [Schapire and Freund, 2012]*
  *Suppose **AdaBoost** is run for T rounds on $m \geq T$ random examples, using base classifiers from a **finite space** $\mathcal{H}$. Then, with probability at least $1 - \delta$, the combined classifier H satisfies*

$$L_{\mathcal{P},c}(H) \leq \widehat{L}_m(H) + \sqrt{\frac{2T\left(\log|\mathcal{H}| + \log(em/T)\right)}{m}} + \sqrt{\frac{\log(1/\delta)}{2m}} \tag{8}$$

  Furthermore, with probability at least $1 - \delta$, if $\mathcal{H}$ is realizable with the training set (i.e. $\widehat{L}_m(h) \equiv 0$), then

$$L_{\mathcal{P},c}(H) \leq \frac{2T\left(\log|\mathcal{H}| + \log(2em/T)\right) + 2\log(2/\delta)}{m}. \tag{9}$$

2.4 Generalization Error Bounds via VC Dimension

- **Lemma 2.7** *(**Growth Number of Combined Hypothesis Class, VC Class**). [Schapire and Freund, 2012]*
  *Assume $\mathcal{H}$ has **finite VC-dimension** $d \geq 1$. Let $m \geq \max\{T, d\} \geq 1$. For any set $\mathcal{D}$ of m points, the number of dichotomies realizable by $L(\mathcal{H}, T)$ is bounded as follows:*

$$|L(\mathcal{H}, T)| \leq \tau_{L(\mathcal{H}, T)}(m) \leq \left(\frac{em}{T}\right)^T \left(\frac{em}{d}\right)^{dT}. \tag{10}$$

7

- **Lemma 2.8** *(VC-Dimension of Combined Hypothesis Class, VC Class).* [Schapire and Freund, 2012, Shalev-Shwartz and Ben-David, 2014]
  Assume $\mathcal{H}$ has **finite VC-dimension** $\nu(\mathcal{H}) = d$ and $\min\{T, d\} \geq 3$. Then the VC dimension of combined hypothesis class is bounded by

$$\nu(L(\mathcal{H}, T)) \leq T(d+1)(3\log(T(d+1)) + 2) = \mathcal{O}(Td\log(Td)). \tag{11}$$

- **Remark** (*Lower Bound on VC Dimension*). [Shalev-Shwartz and Ben-David, 2014]
  For some base hypothesis class $\mathcal{H}$, the VC-dimension of ensemble is at least Td. For instance, for $\mathcal{H}_n$ be the class of *decision stumps* over $\mathbb{R}^n$, we can show that $\log(n) \leq d = \nu(\mathcal{H}) \leq 2\log(n) + 5$. In this example, for all $T \geq 1$,

$$\nu(L(\mathcal{H}_n, T)) \geq 0.5T\log(n) \asymp \Omega(Td).$$

- **Theorem 2.9** *(Generalization Bound for AdaBoost via VC Dimension).* [Schapire and Freund, 2012]
  Suppose **AdaBoost** is run for T rounds on $m \geq \max\{T, d\}$ random examples, using base classifiers from a **finite space** $\mathcal{H}$. Then, with probability at least $1 - \delta$, the combined classifier H satisfies

$$L_{\mathcal{P},c}(H) \leq \widehat{L}_m(H) + \sqrt{\frac{2T(d\log(em/d) + \log(em/T))}{m}} + \sqrt{\frac{\log(1/\delta)}{2m}} \tag{12}$$

Furthermore, with probability at least $1 - \delta$, if $\mathcal{H}$ is realizable with the training set (i.e. $\widehat{L}_m(h) \equiv 0, \forall h \in \mathcal{H}$), then

$$L_{\mathcal{P},c}(H) \leq \frac{2T(d\log(2em/d) + \log(2em/T)) + 2\log(2/\delta)}{m}. \tag{13}$$

- **Corollary 2.10** *[Schapire and Freund, 2012]*
  Assume, in addition to the assumptions of theorem 2.9, that each base classifier has weighted error $\epsilon_t \leq 1/2 - \gamma$ for some $\gamma > 0$. Let the number of rounds T be equal to

$$\inf\left\{t \in \mathbb{N} : t \geq \frac{\log(m)}{2\gamma^2}\right\}$$

Then, with probability at least $1 - \delta$, the generalization error of the combined classifier H will be at most

$$\mathcal{O}\left(\frac{1}{m}\left[\frac{\log(m)}{\gamma^2}\left(\log(m) + d\log\left(\frac{m}{d}\right)\right) + \frac{1}{\delta}\right]\right)$$

- **Remark** Ignoring the log factor, the generalization error bound (12) can be summarized as

$$L_{\mathcal{P},c}(H) \leq \widehat{L}_m(H) + \mathcal{O}\left(\sqrt{\frac{T\mathcal{C}_{\mathcal{H}}}{m}}\right)$$

where $\mathcal{C}_{\mathcal{H}}$ is some complexity measure of base class $\mathcal{H}$.

- **Theorem 2.11** *(Strong Learnable = Weak Learnable)* *[Schapire and Freund, 2012]*
  A target class $\mathcal{H}$ is (efficiently) **weakly** PAC learnable **if and only if** it is (efficiently) **strongly** PAC learnable.

2.5 Generalization Error Bounds via Large Margin Theory

- **Remark** (*Limit of VC Dimension Analysis*)
  The upper bound grows as $\mathcal{O}(dT\log(dT))$, thus the bound suggests that **AdaBoost** *could overfit* for **large values of** T, and indeed this can occur. However, in many cases, it has been observed empirically that *the generalization error of AdaBoost **decreases** as a function of the number of rounds* of boosting T.

- **Definition** (L_1-**Margin**) [Mohri et al., 2018, Schapire and Freund, 2012]
  The L_1-margin $\rho(x)$ of a point $x \in \mathcal{X}$ with label $y \in \{-1, +1\}$ for *a linear combination of base classifiers* $g = \sum_{t=1}^{T} \alpha_t h_t = \langle \alpha\,, h\rangle$ with $\alpha \neq 0$ and $h_t \in \mathcal{H}$ for all $t \in [1, T]$ is defined as

$$\rho(x) := y\frac{\langle \alpha\,, h(x)\rangle}{\|\alpha\|_1} = y\frac{\sum_{t=1}^{T} \alpha_t h_t(x)}{\|\alpha\|_1} \tag{14}$$

  The L_1-margin of a linear combination classifier g **with respect to a sample** $\mathcal{D}$ is **the minimum margin** of the points within the sample:

$$\rho := \min_{i=1,\ldots,m} y_i \frac{\langle \alpha\,, h(x_i)\rangle}{\|\alpha\|_1} = \min_{i=1,\ldots,m} \frac{\sum_{t=1}^{T} \alpha_t y_i h_t(x_i)}{\|\alpha\|_1} \tag{15}$$

- **Remark** When the coefficients α_t are **non-negative**, as in the case of *AdaBoost*, $\rho(x)$ is **a convex combination** of the base classifier values $h_t(x)$. In particular, if the base classifiers h_t take values in $[-1, +1]$, then $\rho(x)$ is in $[-1, +1]$. *The absolute value $|\rho(x)|$ can be interpreted* as **the confidence** *of the classifier g in that label.*

- **Definition** (*Convex Hull of Hypothesis Class*)
  For any hypothesis class $\mathcal{H}$, **the convex hull** of set $\mathcal{H}$, denoted as $\mathrm{conv}(\mathcal{H})$, is defined as

$$\mathrm{conv}(\mathcal{H}) := \left\{\sum_{k=1}^{T} \lambda_k h_k(\cdot) : T \geq 1, \forall k \in [1, T], \lambda_k \geq 0, h_k \in \mathcal{H}, \sum_{k=1}^{T} \lambda_k \leq 1\right\}.$$

- **Definition** (*Empirical Rademacher Complexity*)
  Let $\mathcal{G}$ be a family of functions mapping from $\mathcal{Z} := \mathcal{X} \times \mathcal{Y}$ to $[a, b]$ and $\mathcal{D} = (z_1, \ldots, z_n)$ a fixed *sample* of size n with elements in $\mathcal{Z}$. Then, **the empirical Rademacher complexity** of $\mathcal{G}$ *with respect to the sample $\mathcal{D}$* is defined as:

$$\widehat{\mathfrak{R}}_{\mathcal{D}}(\mathcal{G}) = \mathbb{E}_\sigma\left[\sup_{g \in \mathcal{G}} \frac{1}{n}\sum_{i=1}^{n} \sigma_i g(z_i)\right] \tag{16}$$

  where $\sigma := (\sigma_1, \ldots, \sigma_n)$ are **independent uniform random variables** taking values in $\{-1, +1\}$. The random variables σ_i are called **Rademacher variables**.

- **Proposition 2.12** (*Empirical Rademacher Complexity of a Convex Hull of Function Class*
  Let $\mathcal{H}$ be a set of functions mapping from $\mathcal{X}$ to $\mathbb{R}$. Then, for any sample $\mathcal{D}$, the empirical Rademacher complexity

$$\widehat{\mathfrak{R}}_{\mathcal{D}}(conv(\mathcal{H})) = \widehat{\mathfrak{R}}_{\mathcal{D}}(\mathcal{H}) \tag{17}$$

  *where $conv(\mathcal{H})$ is **the convex hull** of set $\mathcal{H}$.*

3 Fundamental Perspectives

3.1 Game Theory

3.2 Online Learning

3.3 Maximum Entropy Estimation

3.4 Iterative Projection Algorithms and Convergence Analysis

References

Trevor Hastie, Robert Tibshirani, and Jerome H Friedman. *The elements of statistical learning: data mining, inference, and prediction*, volume 2. Springer, 2009.

Mehryar Mohri, Afshin Rostamizadeh, and Ameet Talwalkar. *Foundations of machine learning.* MIT press, 2018.

Robert E. Schapire and Yoav Freund. *Boosting: Foundations and Algorithms.* The MIT Press, 2012. ISBN 0262017180.

Shai Shalev-Shwartz and Shai Ben-David. *Understanding machine learning: From theory to algorithms.* Cambridge university press, 2014.