Lecture 4: Empirical Processes

Tianpei Xie

Feb. 1st., 2023

Contents

1 Uniform Law of Large Numbers

1.1 Motivations

- **Remark** (***Unbiased Estimator of Cumulative Distribution Function***)
  The law of any scalar random variable X can be fully specified by its ***cumulative distribution function (CDF)***, whose value at any point $t \in \mathbb{R}$ is given by $F(t) := \mathcal{P}[X \leq t]$. Now suppose that we are given a collection $\{X_i\}_{i=1}^{n}$ of n i.i.d. samples, each drawn according to the law specified by F. A natural *estimate* of F is ***the empirical CDF*** given by

$$\widehat{F}_n(t) := \frac{1}{n}\sum_{i=1}^{n} \mathbb{1}_{(-\infty,t]}(X_i), \tag{1}$$

  where $\mathbb{1}_{(-\infty,t]}(x)$ is a $\{0,1\}$-valued indicator function for the event $\{x \leq t\}$. Since ***the population CDF*** can be written as $F(t) = \mathbb{E}\left[\mathbb{1}_{(-\infty,t]}(X)\right]$, the empirical CDF is an ***unbiased estimate***.

  For each $t \in \mathbb{R}$, ***the strong law of large numbers*** suggests that

$$\widehat{F}_n(t) \to F(t), \qquad \text{a.s.}$$

  A natural goal is to strengthen *this pointwise convergence* to a form of ***uniform convergence***. The reason why uniform convergence of $\widehat{F}_n(t)$ to $F(t)$ is important is that it can be used to prove the ***consistency*** of ***plug-in estimator*** for *functionals of distribution function.*

- **Example** (***Expectation Functionals***)
  Given some integrable function g, we may define ***the expectation functional*** γ_g via

$$\gamma_g(F) := \int g(x)dF(x). \tag{2}$$

  For any g, *the plug-in estimate* is given by $\gamma_g(\widehat{F}_n) = \frac{1}{n}\sum_{i=1}^{n} g(X_i)$, corresponding to ***the sample mean*** of $g(X)$.

- **Example** (***Quantile Functionals***)
  For any $\alpha \in [0,1]$, ***the quantile functional*** Q_α is given by

$$Q_\alpha(F) := \inf\left\{t \in \mathbb{R} : F(t) \geq \alpha\right\}. \tag{3}$$

  The ***median*** corresponds to the special case $\alpha = 0.5$. *The plug-in estimate* is given by

$$Q_\alpha(\widehat{F}_n) := \inf\left\{t \in \mathbb{R} : \frac{1}{n}\sum_{i=1}^{n} \mathbb{1}_{(-\infty,t]}(X_i) \geq \alpha\right\} \tag{4}$$

  and corresponds to estimating the α-th quantile of the distribution by *the α-th* ***sample quantile***. In the special case $\alpha = 0.5$, this estimate corresponds to *the sample median*. In this case, $Q_\alpha(\widehat{F}_n)$ is a fairly complicated, *nonlinear function of all the variables*, so that this convergence does not follow immediately by a classical result such as the law of large numbers.

- **Example** (***Goodness-of-fit Functionals***)
  It is frequently of interest to test the hypothesis of whether or not a given set of data has

been drawn from a known distribution F_0. Such tests can be performed using *functionals that **measure the distance** between F and the target CDF F_0*, including *the sup-norm distance* $\|F - F_0\|_\infty$, or other distances such as ***the Cramer-von Mises criterion*** based on the functional

$$\gamma_g(F) := \int_{-\infty}^{+\infty} (F(x) - F_0(x))^2 \, dF_0(x)$$

- **Remark** (***Consistency of Plug-In Estimate***)
  For any ***plug-in estimator*** $\gamma_g(\widehat{F}_n)$, an important question is to understand when it is ***consistent*** – that is, when does $\gamma_g(\widehat{F}_n)$ converge to $\gamma_g(F)$ in *probability* (or *almost surely*)?

  We can define *the* ***continuity*** *of a* ***functional*** γ with respect to *the supremum norm*: more precisely, we say that the functional γ is ***continuous*** *at F in the* ***sup-norm*** if, for all $\epsilon > 0$, there exists a $\delta > 0$ such that

  $$\|G - F\|_\infty := \sup_{t \in \mathbb{R}} |G(t) - F(t)| \leq \delta \quad \text{implies that} \quad |\gamma(G) - \gamma(F)| \leq \epsilon.$$

  Thus for any ***continuous functional***, it reduces the ***consistency*** question for *the plug-in estimator* $\gamma_g(\widehat{F}_n)$ to the issue of whether or not the random variable $\|\widehat{F}_n - F\|_\infty$ ***converges to zero***.

1.2 Glivenko-Cantelli Theorem

- **Theorem 1.1** ***(Glivenko-Cantelli Theorem)*** *[Wellner et al., 2013, Wainwright, 2019, Giné and Nickl, 2021]*
  For any distribution, the empirical CDF

  $$\widehat{F}_n(t) := \frac{1}{n} \sum_{i=1}^{n} \mathbb{1}_{(-\infty, t]}(X_i)$$

  is a ***strongly consistent estimator*** *of the population CDF in* ***the uniform norm****, meaning that*

  $$\left\| \widehat{F}_n - F \right\|_\infty := \sup_{t \in \mathbb{R}} \left| \widehat{F}_n(t) - F(t) \right| \to 0, \;\; a.s. \tag{5}$$

- **Remark** (***Uniform Law of Large Numbers***)
  The Glivenko-Cantelli theorem *generalizes* ***the strong law of large numbers*** to *stochastic process*. It confirms that the *convergence* of sample mean $\mathcal{P}_n f$ to its expectation $\mathcal{P} f$ is true in function space $\mathcal{F}$ not only *in* ***pointwise topology*** but also *in* ***uniform topology***. Thus, *the Glivenko-Cantelli theorem* is also called ***the uniform law of large numbers***.

2 Empirical Processes

2.1 Definitions

- **Definition** (***Empirical Measure***) [Wellner et al., 2013, Giné and Nickl, 2021]
  Let $(\mathcal{X}, \mathscr{F}, \mathcal{P})$ be a *probability space*, and let $X_i, i \in \mathbb{N}$, be the *coordinate functions* of *the*

infinite product probability space $(\Omega, \mathscr{B}, \mathbb{P}) := (\mathcal{X}^{\infty}, \mathscr{F}^{\infty}, \mathcal{P}^{\infty})$, $X_i : \mathcal{X}^{\infty} \to \mathcal{X}$, which are ***independent*** *identically distributed* $\mathcal{X}$-valued random variables with law $\mathcal{P}$.

The empirical measure corresponding to the '*observations*' $X_1, \ldots, X_n$, for any $n \in \mathbb{N}$, is defined as ***the random discrete probability measure***

$$\mathcal{P}_n := \frac{1}{n} \sum_{i=1}^{n} \delta_{X_i} \tag{6}$$

where δ_x is *Dirac measure* at x, that is, unit mass at the point x. In other words, for each event A, $\mathcal{P}_n(A)$ is *the* ***proportion*** *of* ***observations*** X_i, $i = 1, \ldots, n$, *that fall in* A; that is,

$$\mathcal{P}(A) = \frac{1}{n} \sum_{i=1}^{n} \delta_{X_i}(A) = \frac{1}{n} \sum_{i=1}^{n} \mathbb{1}\left\{X_i \in A\right\}, \quad A \in \mathscr{F}.$$

- **Remark** (***Probability Measure with Operator Notation***) [Wellner et al., 2013, Giné and Nickl, 2021]
  For any measure μ and μ-integrable function f, we will use the following ***operator notation*** for the integral of f with respect to μ:
  $$\mu f \equiv \mu(f) = \int_{\Omega} f d\mu.$$
  This is valid since there exists an isomorphism between *the space of probability measure* and *the space of bounded linear functional* on $\mathcal{C}_0(\Omega)$ by Riesz-Markov representation theorem (assuming Ω is *locally compact*). By this notion the expectation $\mathcal{P}f = \mathbb{E}_{\mathcal{P}}\left[f\right]$.

- **Definition** (***Empirical Process***) [Wellner et al., 2013, Giné and Nickl, 2021]
  Let $\mathcal{F}$ be a *collection of* $\mathcal{P}$*-integrable functions* $f : \mathcal{X} \to \mathbb{R}$, usually infinite. For any such class of functions $\mathcal{F}$, ***the empirical measure*** defines a ***stochastic process***
  $$f \to \mathcal{P}_n f, \quad f \in \mathcal{F} \tag{7}$$
  which we may call ***the empirical process indexed by*** $\mathcal{F}$, although we prefer to reserve the notation '*empirical process*' for *the* ***centred*** *and* ***normalised*** *process*
  $$f \to \nu_n(f) := \sqrt{n}\left(\mathcal{P}_n f - \mathcal{P} f\right), \quad f \in \mathcal{F}. \tag{8}$$

- **Remark** An explicit notion of *(centered and normalized) empirical process* is
  $$\sqrt{n}\left(\mathcal{P}_n f - \mathcal{P} f\right) \equiv \frac{1}{\sqrt{n}} \sum_{i=1}^{n} \left(f(X_i) - \mathbb{E}_{\mathcal{P}}\left[f(X)\right]\right), \quad f \in \mathcal{F}.$$
  where $X_1, \ldots, X_n \sim \mathcal{P}$ are i.i.d random variables. Note that it is a stochastic process since *the function* f *is changing* in $\mathcal{F}$, i.e. the process $\left(\mathcal{P}_n - \mathcal{P}\right) f$ is indexed by function $f \in \mathcal{F}$ not finite dimensional variable.

- **Remark** (***Random Measure***)
  Normally we assume that data are sampled from some distribution $\mathcal{P}$ and the data itself is random. However, the empirical measure
  $$\mathcal{P}_n := \frac{1}{n} \sum_{i=1}^{n} \delta_{X_i}$$

itself is considered as a ***random*** probability measure. That is, *the sampling mechanism itself contains randomness* and it is not sampling from one distribution but ***a system of distributions depending on the choice of dataset*** $X_1, \ldots, X_n$, which in turn were sampled from some *prior* $\mathcal{P}$. Due to this randomness, $\mathcal{P}_n f = \mathbb{E}_{\mathcal{P}_n}[f]$ is not a fixed expectation number but a random variable. In fact, this is the empirical mean (i.e. sample mean)

$$\mathcal{P}_n f = \mathbb{E}_{\mathcal{P}_n}[f] = \frac{1}{n}\sum_{i=1}^{n} f(X_i).$$

The critical difference between mean of empirical measure vs. sample mean is that we now assume that f is ***not fixed***.

- **Remark** *In probability theory*, an ***empirical process*** is a ***stochastic process*** that *describes the **proportion** of objects in a system **in a given state***. Applications of the theory of empirical processes arise in ***non-parametric statistics***.

- **Remark** (***Object of Empirical Process Theory***)
  *The **object** of empirical process theory* is to study *the **properties** of the **approximation*** of $\mathcal{P}f$ by $\mathcal{P}_n f$, ***uniformly in*** $\mathcal{F}$, concretely, to obtain both ***probability estimates*** *for the random quantities*

  $$\left\|\mathcal{P}_n - \mathcal{P}\right\|_{\mathcal{F}} := \sup_{f\in\mathcal{F}} \left|\mathcal{P}_n f - \mathcal{P}f\right|$$

  and ***probabilistic limit theorems*** for the processes $\{(\mathcal{P}_n - \mathcal{P})(f) : f \in \mathcal{F}\}$.

  Note that the quantity $\left\|\mathcal{P}_n - \mathcal{P}\right\|_{\mathcal{F}}$ is a ***random variable*** since $\mathcal{P}_n$ is a ***random measure***.

- **Remark** (***Measurability Problem***)
  There may be a ***measurability problem*** for

  $$\left\|\mathcal{P}_n - \mathcal{P}\right\|_{\mathcal{F}} := \sup_{f\in\mathcal{F}} \left|\mathcal{P}_n f - \mathcal{P}f\right|$$

  since *the **uncountable** suprema* of measurable functions *may not be measurable.*

  However, there are many situations where this is actually a ***countable supremum***. For instance, for probability distribution on $\mathbb{R}$

  $$\left\|\mathcal{P}_n - \mathcal{P}\right\|_{\infty} := \sup_{t\in\mathbb{R}} \left|(\mathcal{P}_n - \mathcal{P})(-\infty, t)\right| = \sup_{t\in\mathbb{Q}} \left|F_n(t) - F(t)\right| = \sup_{t\in\mathbb{Q}} \left|(\mathcal{P}_n - \mathcal{P})(-\infty, t)\right|$$

  where $F(t) = \mathcal{P}(-\infty, t)$ is the cumulative distribution function. If $\mathcal{F}$ is *countable* or if there exists $\mathcal{F}_0$ *countable* such that

  $$\left\|\mathcal{P}_n - \mathcal{P}\right\|_{\mathcal{F}} = \left\|\mathcal{P}_n - \mathcal{P}\right\|_{\mathcal{F}_0}, \quad \text{a.s.}$$

  then the measurability problem disappears.

  For the next few sections we will simply assume that the class $\mathcal{F}$ is *countable.*

- **Remark** (***Bounded Assumption***)
  If we assume that

  $$\sup_{f\in\mathcal{F}} \left|f(x) - \mathcal{P}f\right| < \infty, \quad \forall x \in \mathcal{X}, \tag{9}$$

then the maps from $\mathcal{F}$ to $\mathbb{R}$,

$$f \to f(x) - \mathcal{P}f, \quad x \in \mathcal{X},$$

are ***bounded functionals*** over $\mathcal{F}$, and therefore, so is $f \to (\mathcal{P}_n - \mathcal{P})(f)$. That is,

$$\mathcal{P}_n - \mathcal{P} \in \ell_\infty(\mathcal{F}),$$

where $\ell_\infty(\mathcal{F})$ is ***the space of bounded real functionals*** *on* $\mathcal{F}$, a *Banach space* if we equip it with the supremum norm $\|\cdot\|_{\mathcal{F}}$.

A large literature is available on *probability in* ***separable Banach spaces***, but unfortunately, $\ell_\infty(\mathcal{F})$ is ***only separable*** when the class $\mathcal{F}$ is ***finite***, and ***measurability problems*** arise because *the probability law* of the process $\{(\mathcal{P}_n - \mathcal{P})(f) : f \in \mathcal{F}\}$ ***does not extend*** *to the Borel σ-algebra of* $\ell_\infty(\mathcal{F})$ even in simple situations.

- **Remark** This chapter addresses ***three main questions*** about the empirical process:

  1. The first question has to do with ***concentration*** of $\|\mathcal{P}_n - \mathcal{P}\|_{\mathcal{F}}$ about *its* ***mean*** when $\mathcal{F}$ is ***uniformly bounded***. Recall that $\|\mathcal{P}_n - \mathcal{P}\|_{\mathcal{F}}$ is a random variable itself, due to randomness of the empirical measure. We mainly use the *non-asymptotic analysis* to obtain *the exponential bound for concentration.*

  2. The second question is do ***good estimates*** for ***mean*** $\mathbb{E}\left[\|\mathcal{P}_n - \mathcal{P}\|_{\mathcal{F}}\right]$ exist? We will examine two main techniques that give answers to this question, both related to ***metric entropy*** and ***chaining***. One of them, called ***bracketing***, uses *chaining* in combination with *truncation* and *Bernstein's inequality*. The other one applies to ***Vapnik-Cervonenkis (VC) classes of functions***.

  3. Finally, the last question about the empirical process refers to ***limit theorems***, mainly ***the uniform law of large numbers*** and the ***central limit theorem***, in fact, the analogues of ***the classical Glivenko-Cantelli*** and *Donsker theorems* for the empirical distribution function.

     Formulation of *the central limit theorem* will require some more *measurability* because we will be considering ***convergence in law*** of random elements in ***not necessarily separable Banach spaces***.

2.2 Glivenko-Cantelli Class

- **Definition** (***Glivenko-Cantelli Class***) [Wellner et al., 2013, Wainwright, 2019, Giné and Nickl, 2021]
  We say that $\mathcal{F}$ is a ***Glivenko-Cantelli class*** for $\mathcal{P}$ if

  $$\|\mathcal{P}_n - \mathcal{P}\|_{\mathcal{F}} := \sup_{f \in \mathcal{F}} |\mathcal{P}_n f - \mathcal{P}f| \to 0$$

  in probability as $n \to \infty$.

  This notion can also be defined in a *stronger* sense, requiring ***almost sure convergence*** of $\|\mathcal{P}_n - \mathcal{P}\|_{\mathcal{F}}$, in which case we say that $\mathcal{F}$ satisfies a ***strong Glivenko-Cantelli law***.

- **Example** (***Empirical CDFs and Indicator Functions***)
Consider the function class

$$\mathcal{F} := \left\{\mathbb{1}_{(-\infty,t]}(\cdot), t \in \mathbb{R}\right\} \tag{10}$$

where $\mathbb{1}_{(-\infty,t]}$ is the $\{0,1\}$-*valued indicator function* of the interval $(-\infty, t]$. For each fixed $t \in \mathbb{R}$, we have the equality $\mathbb{E}\left[\mathbb{1}_{(-\infty,t]}(X)\right] = \mathcal{P}[X \le t] = F(t)$, so that the classical *Glivenko-Cantelli theorem* is equivalent to a ***strong uniform law*** *for the class* (10),

2.3 Tail bounds for Empirical Processes

- **Remark** Consider *the* ***suprema*** *of* ***empirical process***:

$$Z := \sup_{f\in\mathcal{F}} \{\mathcal{P}_n f\} = \sup_{f\in\mathcal{F}} \left\{\frac{1}{n}\sum_{i=1}^{n} f(X_i)\right\} \tag{11}$$

where $(X_1, \ldots, X_n)$ are independent random variables drawn from $\mathcal{P} := \otimes_{i=1}^{n}\mathcal{P}_i$, each $\mathcal{P}_i$ is supported on some set $\mathcal{X}_i \subseteq \mathcal{X}$. $\mathcal{F}$ is a family of real-valued functions $f : \mathcal{X} \to \mathbb{R}$. The primary goal of this section is to derive a number of *upper bounds* on *the tail event* $\{Z \ge \mathbb{E}[Z] + t\}$.

- **Theorem 2.1** ***(Functional Hoeffding Inequality)*** *[Wainwright, 2019, Boucheron et al., 2013]*
For each $f \in \mathcal{F}$ and $i = 1, \ldots, n$, assume that there are real numbers $a_{i,f} \le b_{i,f}$ such that $f(x) \in [a_{i,f}, b_{i,f}]$ for all $x \in \mathcal{X}_i$. Then for all $t \ge 0$, we have

$$\mathcal{P}\{Z \ge \mathbb{E}[Z] + t\} \le \exp\left(-\frac{nt^2}{4L^2}\right) \tag{12}$$

where $Z := \sup_{f\in\mathcal{F}}\left\{\frac{1}{n}\sum_{i=1}^{n} f(X_i)\right\}$, and $L^2 := \sup_{f\in\mathcal{F}}\left\{\frac{1}{n}\sum_{i=1}^{n}(a_{i,f} - b_{i,f})^2\right\}$.

- **Theorem 2.2** ***(Functional Bernstein Inequality, Talagrand Concentration for Empirical Processes)*** *[Wainwright, 2019, Boucheron et al., 2013]*
Consider a ***countable*** *class of functions $\mathcal{F}$* ***uniformly bounded*** *by b. Then for all $t > 0$, the suprema of empirical process Z as defined in (11) satisfies the upper tail bound*

$$\mathcal{P}\{Z \ge \mathbb{E}[Z] + t\} \le \exp\left(-\frac{nt^2}{8e\Sigma^2 + 4bt}\right) \tag{13}$$

where $\Sigma^2 := \mathbb{E}\left[\sup_{f\in\mathcal{F}}\left\{\frac{1}{n}\sum_{i=1}^{n} f^2(X_i)\right\}\right]$ is ***the weak variance.***

- **Remark** As opposed to control only in terms of ***bounds*** *on the* ***function values***, the inequality (13) ***also*** brings a notion of ***variance*** into play.

- **Remark** We will prove the bound in next section:

$$\Sigma^2 \le \sigma^2 + 2b\mathbb{E}[Z]$$

where $\sigma^2 := \sup_{f\in\mathcal{F}} \mathbb{E}\left[f^2(X)\right]$. Then, the functional Bernstein inequality (13) can be formulated as

$$\mathcal{P}\left\{Z \ge \mathbb{E}[Z] + c_0\gamma\sqrt{t} + c_1 bt\right\} \le e^{-nt} \tag{14}$$

for some constant c_0, c_1 and $\gamma^2 := \sigma^2 + 2b\mathbb{E}[Z]$. We can have an alternative form of this bound (14) for any $\epsilon > 0$,

$$\mathcal{P}\left\{Z \geq (1+\epsilon)\mathbb{E}[Z] + c_0\sigma\sqrt{t} + (c_1 + c_0^2/\epsilon)bt\right\} \leq e^{-nt}. \tag{15}$$

- **Theorem 2.3** ***(Bousquet's Inequality, Functional Bennet Inequality)*** *[Boucheron et al., 2013]*
  Let $X_1, \ldots, X_n$ be ***independent identically distributed*** *random vectors. Assume that $\mathbb{E}[f(X_i)] = 0$, and that $f(X_i) \leq 1$ for all $f \in \mathcal{F}$. Let*

$$\gamma^2 = \sigma^2 + 2\mathbb{E}[Z],$$

  where $\sigma^2 := \sup_{f\in\mathcal{F}}\left\{\frac{1}{n}\sum_{i=1}^n \mathbb{E}\left[f^2(X_i)\right]\right\}$ is ***the wimpy variance****. Let $\phi(u) = e^u - u - 1$ and $h(u) = (1+u)\log(1+u) - u$, for $u \geq -1$. Then for all $\lambda \geq 0$,*

$$\log \mathbb{E}\left[e^{\lambda(Z-\mathbb{E}[Z])}\right] \leq n\gamma^2\phi(\lambda).$$

  Also, for all $t \geq 0$,

$$\mathcal{P}\left\{Z \geq \mathbb{E}[Z] + t\right\} \leq \exp\left(-n\gamma^2\, h\left(\frac{t}{\gamma^2}\right)\right). \tag{16}$$

2.4 Maximal Inequalities

3 Variance of Suprema of Empirical Process

3.1 General Upper Bounds for the Variance

- **Definition** (***Variances of Empirical Process***)
  Let $X_1, \ldots, X_n$ be independent random variables taking values in $\mathcal{X}$. Depending on ***ordering*** of the ***expectation***, ***suprema*** and ***summation*** operator, we define *three different types of* ***variance*** associated with empirical process

$$\mathcal{P}_n f = \mathbb{E}_{\mathcal{P}_n}[f] = \frac{1}{n}\sum_{i=1}^n f(X_i).$$

  1. ***The strong variance*** is defined as

$$V := \frac{1}{n}\sum_{i=1}^n \mathbb{E}\left[\sup_{f\in\mathcal{F}} f^2(X_i)\right] \tag{17}$$

  2. ***The weak variance*** is defined as

$$\Sigma^2 := \mathbb{E}\left[\sup_{f\in\mathcal{F}}\left\{\frac{1}{n}\sum_{i=1}^n f^2(X_i)\right\}\right] \tag{18}$$

3. ***The wimpy variance*** is defined as

$$\sigma^2 := \sup_{f \in \mathcal{F}} \left\{ \frac{1}{n} \sum_{i=1}^{n} \mathbb{E}\left[f^2(X_i) \right] \right\} \tag{19}$$

By Jensen's inequality,

$$\sigma^2 \leq \Sigma^2 \leq V$$

In general, there may be *significant gaps between any two of these quantities*. A notable difference is the case of ***Rademacher averages*** when $\sigma^2 = \Sigma^2$.

3.2 Symmetrization and Contraction Principle

3.3 Bounding the Weak Variance via Wimpy Variance

3.4 Rademacher Complexity and Gaussian Complexity

4 Expected Value of Suprema of Empirical Process

4.1 Covering Number, Packing Number and Metric Entropy

4.2 Chaining and Dudley's Entropy Integral

4.3 Vapnik-Chervonenkis Class

4.4 Comparison Theorems

References

Stéphane Boucheron, Gábor Lugosi, and Pascal Massart. *Concentration inequalities: A nonasymptotic theory of independence*. Oxford university press, 2013.

Evarist Giné and Richard Nickl. *Mathematical foundations of infinite-dimensional statistical models*. Cambridge university press, 2021.

Martin J Wainwright. *High-dimensional statistics: A non-asymptotic viewpoint*, volume 48. Cambridge University Press, 2019.

Jon Wellner et al. *Weak convergence and empirical processes: with applications to statistics*. Springer Science & Business Media, 2013.